

Running Hadoop On Ubuntu Linux (Single-Node Cluster)

Expt No: 2

March 07, 2019

Author: Subalakshmi Shanthosi S (186001008)

Aim

To configure and install Hadoop in Ubuntu 16.04 LTS OS flavour.

Software's Used

- Ubuntu 16.04 LTS
- Hadoop 2.7.3

Description

Installation of Oracle VirtualBox with guest Operating System as Ubuntu 16.04. Installation of necessary packages namely - openssh-server, openssh-client, java, hadoop in the created virtualbox instance.

Procedure

1. Launch Ubuntu 16.04 LTS.
2. Login to the OS with sudo permission and install the following packages using apt-get command
 - openssh-server
 - openssh-client
 - java jdk 8
 - javac compiler
 - hadoop 2.7.3

Output

```
admin@ssn-c6:~$ sudo apt-get install openssh-server
[sudo] password for admin:
Reading package lists... Done
Building dependency tree
Reading state information... Done
openssh-server is already the newest version (1:7.2p2-4ubuntu2.7).
The following packages were automatically installed and are no longer required:
  libdbusmenu-gtk4 libllvm3.8 libqmt-glib1
Use 'sudo apt autoremove' to remove them.
0 upgraded, 0 newly installed, 0 to remove and 4 not upgraded.
admin@ssn-c6:~$ sudo apt-get install openssh-client
Reading package lists... Done
Building dependency tree
Reading state information... Done
openssh-client is already the newest version (1:7.2p2-4ubuntu2.7).
openssh-client set to manually installed.
The following packages were automatically installed and are no longer required:
  libdbusmenu-gtk4 libllvm3.8 libqmt-glib1
Use 'sudo apt autoremove' to remove them.
0 upgraded, 0 newly installed, 0 to remove and 4 not upgraded.
admin@ssn-c6:~$ sudo apt-get install vi
admin@ssn-c6:~$ sudo apt-get install virtu
```

Figure 1: Install openssh-server,openssh-client in Ubuntu OS.

```
hduser@client-VirtualBox: ~/jdk1.8.0_171/bin
hduser@client-VirtualBox:~/jdk1.8.0_171/bin$ export JAVA_HOME=/home/hduser/jdk1.
3.0_171/bin
hduser@client-VirtualBox:~/jdk1.8.0_171/bin$ export PATH=$PATH:$JAVA_HOME
hduser@client-VirtualBox:~/jdk1.8.0_171/bin$ echo $PATH
/usr/local/bin:/usr/bin:/bin:/usr/local/games:/usr/games:/snap/bin:/home/hduser/
jdk1.8.0_171/bin:/home/hduser/jdk1.8.0_171/bin
hduser@client-VirtualBox:~/jdk1.8.0_171/bin$
```

Figure 2: Setting Java Home environment variable to the specified download path of JDK-1.7.

```
client@client-VirtualBox:~$ sudo addgroup hadoop
Adding group 'hadoop' (GID 1001) ...
Done.
client@client-VirtualBox:~$ sudo adduser --ingroup hadoop hduser
Adding user 'hduser' ...
Adding new user 'hduser' (1001) with group 'hadoop' ...
Creating home directory '/home/hduser' ...
Copying files from '/etc/skel' ...
Enter new UNIX password:
Retype new UNIX password:
passwd: password updated successfully
Changing the user information for hduser
Enter the new value, or press ENTER for the default
  Full Name []:
    Room Number []:
    Work Phone []:
    Home Phone []:
    other []:
Is the information correct? [Y/n]
client@client-VirtualBox:~$
client@client-VirtualBox:~$ su - hduser
```

Figure 3: Adding a dedicated hadoop system user.

```
hduser@client-VirtualBox: ~  
hduser@client-VirtualBox:~$ ssh-keygen -t rsa -P ""  
Generating public/private rsa key pair.  
Enter file in which to save the key (/home/hduser/.ssh/id_rsa):  
Created directory '/home/hduser/.ssh'.  
Your identification has been saved in /home/hduser/.ssh/id_rsa.  
Your public key has been saved in /home/hduser/.ssh/id_rsa.pub.  
The key fingerprint is:  
SHA256:vnkUKXUrfJx9B/DpAEj2uW29mZTH0np5C0m09uGt00o hduser@client-VirtualBox  
The key's randomart image is:  
+---[RSA 2048]-----+  
| .000+... |  
| ..=00.... |  
| 000.00. |  
| .+0.B 0 |  
| S..0+ @ |  
| .. = B |  
| .. . E.. |  
| O. .+..+ |  
| O. ...*+ |  
+----[SHA256]-----+  
hduser@client-VirtualBox:~$
```

Figure 4: Configuring SSH in newly created user.

```
hduser@client-VirtualBox: ~  
hduser@client-VirtualBox:~$ cat /proc/sys/net/ipv6/conf/all/disable_ipv6  
0  
hduser@client-VirtualBox:~$
```

Figure 5: Disabling IPv6 in the newly created user account.

```
# Log Martian Packets  
#net.ipv4.conf.all.log_martians = 1  
  
# /etc/sysctl.conf  
# disable ipv6  
net.ipv6.conf.all.disable_ipv6 = 1  
net.ipv6.conf.default.disable_ipv6 = 1  
net.ipv6.conf.lo.disable_ipv6 = 1
```

Figure 6: Disabling IPv6 in the newly created user account.

```
hduser@client-VirtualBox: ~  
hduser@client-VirtualBox:~$ wget https://archive.apache.org/dist/hadoop/core/hadoop-2.7.3/hadoop-2.7.3.tar.gz  
--2019-03-06 14:33:44-- https://archive.apache.org/dist/hadoop/core/hadoop-2.7.3/hadoop-2.7.3.tar.gz  
Resolving proxy.ssn.net (proxy.ssn.net)... 192.168.2.5  
Connecting to proxy.ssn.net (proxy.ssn.net)[192.168.2.5]:8080... connected.  
Proxy request sent, awaiting response... 200 OK  
Length: 214092195 (204M) [application/x-gzip]  
Saving to: 'hadoop-2.7.3.tar.gz'  
  
hadoop-2.7.3.tar.gz 17%[==>] 35.01M 283KB/s eta 13m 14s
```

Figure 7: Installation of Hadoop 2.7.3 in new user login.

```
<configuration>  
<!-- conf/core-site.xml -->  
<property>  
  <name>hadoop.tmp.dir</name>  
  <value>/app/hadoop/tmp</value>  
  <description>A base for other temporary directories.</description>  
</property>  
  
<property>  
  <name>fs.default.name</name>  
  <value>hdfs://localhost:54310</value>  
  <description>The name of the default file system. A URI whose  
scheme and authority determine the FileSystem implementation. The  
uri's scheme determines the config property (fs.SCHEME.impl) naming  
the FileSystem implementation class. The uri's authority is used  
to  
determine the host, port, etc. for a filesystem.</description>  
</property>  
|  
</configuration>
```

Figure 8: Configuring hadoop core-site.xml .

```
<configuration>  
<!-- conf/mapred-site.xml -->  
<property>  
  <name>mapred.job.tracker</name>  
  <value>localhost:54311</value>  
  <description>The host and port that the MapReduce job tracker runs  
at. If "local", then jobs are run in-process as a single map  
and reduce task.  
</description>  
</property>  
  
</configuration>
```


Figure 9: Configuring Hadoop MapReduce.

```

<configuration>
<!-- conf/hdfs-site.xml -->
<property>
  <name>dfs.replication</name>
  <value>1</value>
  <description>Default block replication.
    The actual number of replications can be specified when the file
    is created.
    The default is used if replication is not specified in create time.
  </description>
</property>
</configuration>

```

Figure 10: Configuring Hadoop HDFS Site.



```

hduser@client-VirtualBox:~$ /usr/local/hadoop-2.7.3/bin/hadoop namenode -format
DEPRECATED: Use of this script to execute hdfs command is deprecated.
Instead use the hdfs command for it.

19/03/13 15:13:57 INFO namenode.NameNode: STARTUP_MSG:
/*****
STARTUP_MSG: Starting NameNode
STARTUP_MSG: host = client-VirtualBox/127.0.1.1
STARTUP_MSG: args = [-format]
STARTUP_MSG: version = 2.7.3
STARTUP_MSG: classpath = /usr/local/hadoop-2.7.3/etc/hadoop:/usr/local/hadoop-2.7.3/share/
hadoop/common/lib/jetty-6.1.26.jar:/usr/local/hadoop-2.7.3/share/hadoop/common/lib/guava-11.
0.2.jar:/usr/local/hadoop-2.7.3/share/hadoop/common/lib/commons-net-3.1.jar:/usr/local/hadoo
p-2.7.3/share/hadoop/common/lib/commons-math3-3.1.1.jar:/usr/local/hadoop-2.7.3/share/hadoop
/common/lib/hadoop-annotations-2.7.3.jar:/usr/local/hadoop-2.7.3/share/hadoop/common/lib/com
mons-io-2.4.jar:/usr/local/hadoop-2.7.3/share/hadoop/common/lib/netty-3.6.2.Final.jar:/usr/l
ocal/hadoop-2.7.3/share/hadoop/common/lib/commons-beanutils-1.7.0.jar:/usr/local/hadoop-2.7.
3/share/hadoop/common/lib/java-xmlbuilder-0.4.jar:/usr/local/hadoop-2.7.3/share/hadoop/commo
n/lib/commons-logging-1.1.3.jar:/usr/local/hadoop-2.7.3/share/hadoop/common/lib/htrace-core-
3.1.0-incubating.jar:/usr/local/hadoop-2.7.3/share/hadoop/common/lib/zookeeper-3.4.6.jar:/us
r/local/hadoop-2.7.3/share/hadoop/common/lib/log4j-1.2.17.jar:/usr/local/hadoop-2.7.3/share/
hadoop/common/lib/curator-framework-2.7.1.jar:/usr/local/hadoop-2.7.3/share/hadoop/common/li
b/api-util-1.0.0-M20.jar:/usr/local/hadoop-2.7.3/share/hadoop/common/lib/commons-beanutils-c
ore-1.8.0.jar:/usr/local/hadoop-2.7.3/share/hadoop/common/lib/servlet-api-2.5.jar:/usr/local
/hadoop-2.7.3/share/hadoop/common/lib/slf4j-api-1.7.10.jar:/usr/local/hadoop-2.7.3/share/had
oop/common/lib/protobuf-java-2.5.0.jar:/usr/local/hadoop-2.7.3/share/hadoop/common/lib/commo
ns-codec-1.4.jar:/usr/local/hadoop-2.7.3/share/hadoop/common/lib/jersey-core-1.9.jar:/usr/lo
cal/hadoop-2.7.3/share/hadoop/common/lib/jersey-iron-1.9.jar:/usr/local/hadoop-2.7.3/share/h

```

Figure 11: Formatting HDFS file system via the NameNode.

Result

Thus the hadoop single node cluster is successfully created in Ubuntu 16.04 OS version and required packages are installed.